Machine Learning Models, Optimisation and Evaluation

1. Introduction

Machine learning (ML) is a subfield of artificial intelligence focused on building systems that learn patterns from data and improve performance on a task without being explicitly programmed. This document provides a structured overview of the main types of machine learning models, model optimisation techniques, and metrics used to assess model performance.

2. Types of Machine Learning Models

2.1. Supervised Learning

Supervised learning models are trained on labelled data, where both inputs and target outputs are known.

2.1.1. Regression Models

Regression models are used when the target variable is continuous. They aim to model the relationship between input features and a numerical output.

- Linear Regression
- Polynomial Regression
- Ridge and Lasso Regression
- Support Vector Regression (SVR)
- Decision Tree Regressors
- Random Forest Regressors
- Gradient Boosting Regressors

2.1.2. Classification Models

Classification models are used when the target variable is categorical. They assign input data to one of several predefined classes.

- Logistic Regression
- k-Nearest Neighbours (k-NN)
- Naive Bayes
- Support Vector Machines (SVM)
- Decision Trees
- Random Forests
- Gradient Boosting (XGBoost, LightGBM, CatBoost)
- Neural Networks

2.2. Unsupervised Learning

Unsupervised learning models operate on unlabelled data to discover hidden structure or patterns.

2.2.1. Clustering

Clustering algorithms group similar data points together based on a distance or similarity measure.

- k-Means
- Hierarchical Clustering
- DBSCAN
- HDBSCAN

- Gaussian Mixture Models (GMM)

2.2.2. Dimensionality Reduction

Dimensionality reduction techniques reduce the number of input features while preserving as much information as possible.
- Principal Component Analysis (PCA)
- t-SNE
- UMAP

2.3. Semi-Supervised Learning

Semi-supervised learning combines a small amount of labelled data with a large amount of unlabelled data to improve learning performance.
- Label Propagation
- Self-training methods

2.4. Reinforcement Learning

Reinforcement learning models learn optimal behaviour through interaction with an environment using reward signals.

2.5. Ensemble Methods

Ensemble methods combine multiple models to improve predictive performance and robustness.
- Bagging (e.g. Random Forests)
- Boosting (AdaBoost, Gradient Boosting)
- Stacking

3. Model Optimisation Techniques

3.1. Feature Engineering

Feature engineering focuses on improving input representations to enhance model performance.
- Feature scaling
  - standardisation
  - normalisation
- Encoding categorical variables
- Feature selection
- Feature extraction

3.2. Optimisation Algorithms

Optimisation algorithms minimise a loss function during training.
- Gradient Descent
- Stochastic Gradient Descent (SGD)
- Mini-batch Gradient Descent
- Momentum
- RMSProp
- Adam

3.3. Hyperparameter Optimisation

Hyperparameter optimisation searches for the best model configuration.

- Grid Search
- Random Search
- Bayesian Optimisation
- Hyperband

3.4. Regularisation Techniques

Regularisation techniques reduce overfitting by constraining model complexity.
- L1 regularisation (Lasso)
- L2 regularisation (Ridge)
- Elastic Net
- Dropout (neural networks)
- Early stopping

3.5. Cross-Validation

Cross-validation estimates a model's ability to generalise to unseen data.
- k-Fold Cross-Validation
- Stratified k-Fold

4. Model Evaluation Metrics

4.1. Regression Metrics

Regression metrics evaluate performance on continuous targets:

- Mean Absolute Error :

$$\mathrm{MAE} = \frac{1}{n}\sum_{i=1}^{n}|y_i - \hat{y}_i|$$

- Mean Squared Error :

$$\mathrm{MSE} = \frac{1}{n}\sum_{i=1}^{n}(y_i - \hat{y}_i)^2$$

- Root Mean Squared Error :

$$\mathrm{RMSE} = \sqrt{\frac{1}{n}\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}$$

- R^2 score : useful to assess the proportion of variance explained by the model, to be used alongside above-mentioned error-based metrics as it does not measure prediction error magnitude.

$$R^2 = 1 - \frac{\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}{\sum_{i=1}^{n}(y_i - \bar{y}_i)^2}$$

4.2. Classification Metrics

Classification metrics assess predictive performance on categorical targets.

- Accuracy : proportion of correct predictions out of all predictions.

$$\mathrm{Accuracy} = \frac{\mathrm{TP} + \mathrm{TN}}{\mathrm{TP} + \mathrm{TN} + \mathrm{FP} + \mathrm{FN}}$$

- Precision
- Recall (true positive rate)
- F1-score

- Confusion matrix

4.2.1. Threshold-Independent Metrics

- Receiver Operating Characteristic (ROC) curve and AUC
- Precision–Recall curve

4.3. Clustering Metrics

Clustering metrics evaluate the quality of unsupervised groupings.
- Silhouette score
- Davies–Bouldin index

4.4. Probabilistic and Ranking Metrics

These metrics are used when models output probabilities or rankings.
- Log loss (cross-entropy) :

$$L = -\sum_{i=1}^{C} (y_i \log(\hat{y}_i))$$

- Brier score
- Mean Average Precision (MAP)
- Normalised Discounted Cumulative Gain (NDCG)

5. Bias–Variance Trade-off

As model complexity increases, bias generally decreases while variance increases. A well-performing model balances both to minimise generalisation error.

6. Conclusion

Machine learning comprises a broad set of models, optimisation strategies, and evaluation metrics. Selecting appropriate techniques depends on the task, data properties, and performance requirements. A structured understanding of these components is essential for developing reliable and effective machine learning systems.